

Breast Cancer Detection

Mt. SAC CISB 62 Midterm Project Fall 2023

Vedavit Shetty

The main objective of this project is to find the best hyperparameters for a neural network model that classifies breast cancer tumors, so that it achieves the highest possible accuracy on the validation set. The Keras Tuner's RandomSearch method is utilized to automate and optimize this hyperparameter tuning process.

You can find this projected hosted on github: <https://github.com/vedavitshetty/CISB-62-Midterm-Breast-Cancer-Detection-Project> (<https://github.com/vedavitshetty/CISB-62-Midterm-Breast-Cancer-Detection-Project>)

Import Libraries

```
In [1]: # Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

Exploratory Data Analysis (EDA)

Load Data

```
In [2]: # Load the dataset
data = load_breast_cancer()
df = pd.DataFrame(data.data, columns=data.feature_names)
df['target'] = data.target
```

Display the first 5 values

```
In [3]: df.head()
```

Out [3]:

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | mean concave points | mean symmetry | |
|---|----------------|-----------------|-------------------|--------------|--------------------|---------------------|-------------------|---------------------------|------------------|--|
| 0 | 17.99 | 10.38 | 122.80 | 1001.0 | 0.11840 | 0.27760 | 0.3001 | 0.14710 | 0.2419 | |
| 1 | 20.57 | 17.77 | 132.90 | 1326.0 | 0.08474 | 0.07864 | 0.0869 | 0.07017 | 0.1812 | |
| 2 | 19.69 | 21.25 | 130.00 | 1203.0 | 0.10960 | 0.15990 | 0.1974 | 0.12790 | 0.2069 | |
| 3 | 11.42 | 20.38 | 77.58 | 386.1 | 0.14250 | 0.28390 | 0.2414 | 0.10520 | 0.2597 | |
| 4 | 20.29 | 14.34 | 135.10 | 1297.0 | 0.10030 | 0.13280 | 0.1980 | 0.10430 | 0.1809 | |

5 rows × 31 columns

See info

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   mean radius                           569 non-null    float64
1   mean texture                           569 non-null    float64
2   mean perimeter                         569 non-null    float64
3   mean area                             569 non-null    float64
4   mean smoothness                        569 non-null    float64
5   mean compactness                       569 non-null    float64
6   mean concavity                          569 non-null    float64
7   mean concave points                    569 non-null    float64
8   mean symmetry                          569 non-null    float64
9   mean fractal dimension                 569 non-null    float64
10  radius error                           569 non-null    float64
11  texture error                           569 non-null    float64
12  perimeter error                         569 non-null    float64
13  area error                             569 non-null    float64
14  smoothness error                       569 non-null    float64
15  compactness error                      569 non-null    float64
16  concavity error                        569 non-null    float64
17  concave points error                   569 non-null    float64
18  symmetry error                         569 non-null    float64
19  fractal dimension error                 569 non-null    float64
20  worst radius                           569 non-null    float64
21  worst texture                           569 non-null    float64
22  worst perimeter                         569 non-null    float64
23  worst area                             569 non-null    float64
24  worst smoothness                       569 non-null    float64
25  worst compactness                      569 non-null    float64
26  worst concavity                        569 non-null    float64
27  worst concave points                    569 non-null    float64
28  worst symmetry                          569 non-null    float64
29  worst fractal dimension                 569 non-null    float64
30  target                                 569 non-null    int64
dtypes: float64(30), int64(1)
memory usage: 137.9 KB
```

See the count, mean, standard deviation, minimum, first quartile, median, third quartile, and maximum values of each column

In [5]:

df.describe()

Out [5]:

| | mean radius | mean texture | mean perimeter | mean area | mean smoothness | mean compactness | mean concavity | |
|-------|----------------|-----------------|-------------------|-------------|--------------------|---------------------|-------------------|---|
| count | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 569.000000 | 5 |
| mean | 14.127292 | 19.289649 | 91.969033 | 654.889104 | 0.096360 | 0.104341 | 0.088799 | |
| std | 3.524049 | 4.301036 | 24.298981 | 351.914129 | 0.014064 | 0.052813 | 0.079720 | |
| min | 6.981000 | 9.710000 | 43.790000 | 143.500000 | 0.052630 | 0.019380 | 0.000000 | |
| 25% | 11.700000 | 16.170000 | 75.170000 | 420.300000 | 0.086370 | 0.064920 | 0.029560 | |
| 50% | 13.370000 | 18.840000 | 86.240000 | 551.100000 | 0.095870 | 0.092630 | 0.061540 | |
| 75% | 15.780000 | 21.800000 | 104.100000 | 782.700000 | 0.105300 | 0.130400 | 0.130700 | |
| max | 28.110000 | 39.280000 | 188.500000 | 2501.000000 | 0.163400 | 0.345400 | 0.426800 | |

8 rows × 31 columns

Check null values

```
In [6]: df.isnull().sum()
```

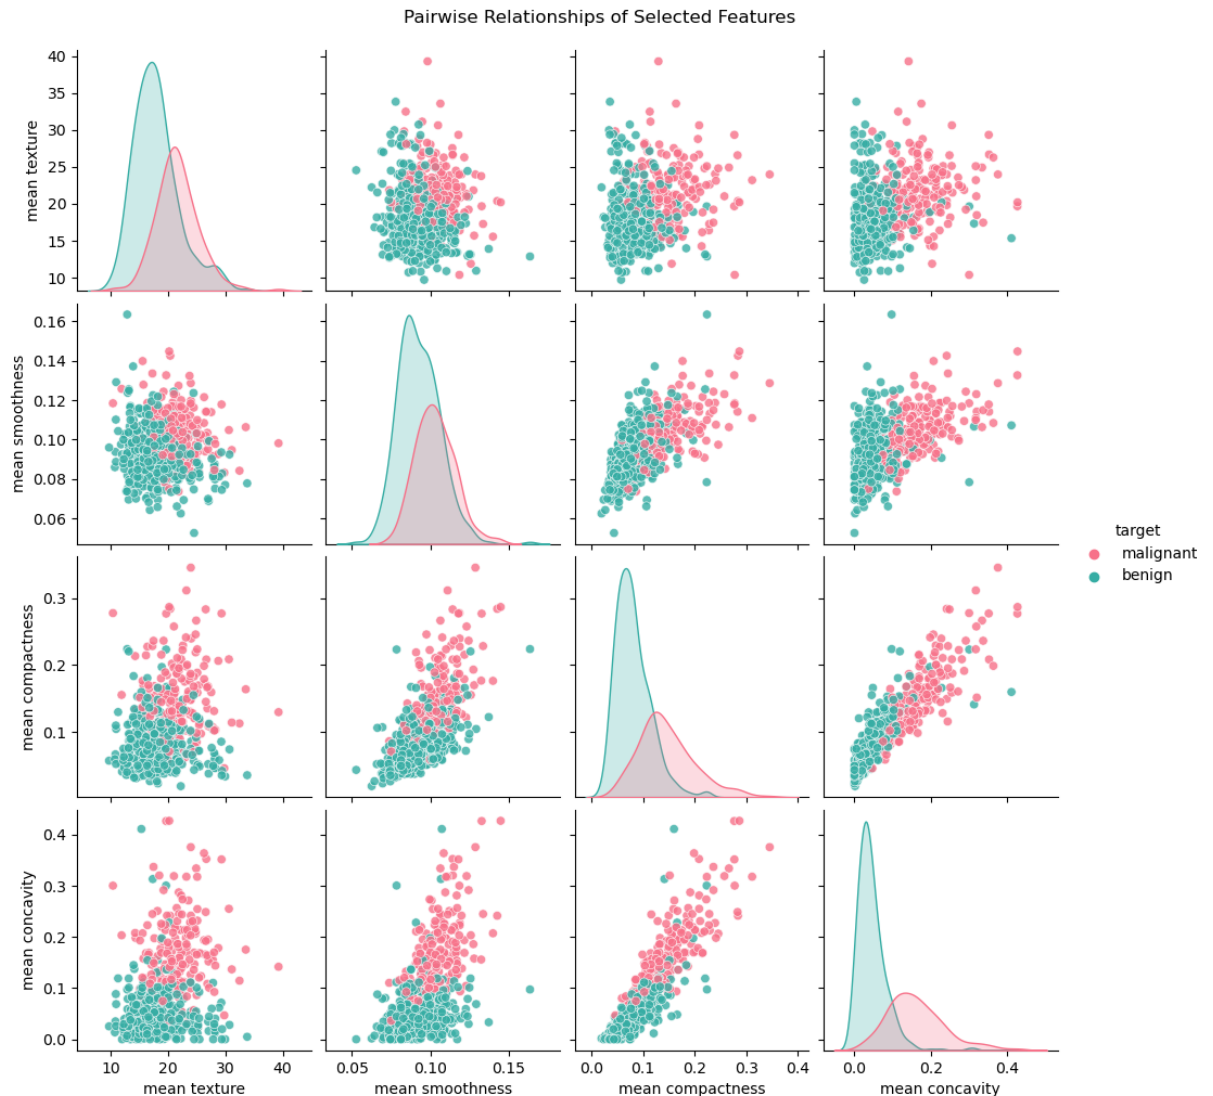
```
Out[6]: mean radius      0
mean texture      0
mean perimeter    0
mean area         0
mean smoothness   0
mean compactness  0
mean concavity    0
mean concave points 0
mean symmetry     0
mean fractal dimension 0
radius error      0
texture error     0
perimeter error   0
area error        0
smoothness error  0
compactness error 0
concavity error   0
concave points error 0
symmetry error    0
fractal dimension error 0
worst radius      0
worst texture     0
worst perimeter   0
worst area        0
worst smoothness  0
worst compactness 0
worst concavity   0
worst concave points 0
worst symmetry    0
worst fractal dimension 0
target           0
dtype: int64
```

Visualization

```
In [7]: selected_features = ['mean texture', 'mean smoothness', 'mean compactness', 'mean concavity']

# Create a temporary dataframe with the target labels
temp_df = df.copy()
temp_df.replace(to_replace='target', {0: data.target_names[0]}), inplace=True
temp_df.replace(to_replace='target', {1: data.target_names[1]}), inplace=True

sns.pairplot(temp_df, hue='target', vars=selected_features, palette="husl")
plt.suptitle('Pairwise Relationships of Selected Features', y=1.02)
plt.show()
```



Insights

- Malignant tumors, in general, tend to have higher values for the features mean smoothness, mean compactness, and mean concavity.
- As mean compactness increases, mean concavity also seems to increase resulting in a strong positive correlation
- Mean smoothness and mean compactness have a more mild positive correlation
- Mean smoothness and mean concavity have some positive correlation, but not as pronounced as the previous 2 mentioned

Data Transformation and Splitting

```
In [8]: # Splitting the data
X = df.drop('target', axis=1)
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,

# Standardizing the data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Build the Deep Learning Model

```
In [9]: import tensorflow as tf
import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout, BatchNormalization

# Building the ANN
model = Sequential()
model.add(Dense(32, activation='relu', input_dim=X_train.shape[1]))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(16, activation='relu'))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['acc
```

Hyperparameter Tuning using Keras Tuner

```
In [11]: #!/pip install keras-tuner

from keras_tuner.tuners import RandomSearch

def build_model(hp):
    model = Sequential()
    model.add(Dense(units=hp.Int('units', min_value=32, max_value=512, step=16)))
    model.add(BatchNormalization())
    model.add(Dropout(0.5))
    model.add(Dense(16, activation='relu'))
    model.add(Dense(1, activation='sigmoid'))
    model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
    return model

tuner = RandomSearch(
    build_model,
    objective='val_accuracy',
    max_trials=5,
    executions_per_trial=2,
    directory='breast_cancer_model_dir',
    project_name='breast_cancer')

tuner.search(X_train, y_train, epochs=10, validation_data=(X_test, y_test))
```

Trial 5 Complete [00h 00m 02s]
val_accuracy: 0.9780701696872711

Best val_accuracy So Far: 0.9780701696872711
Total elapsed time: 00h 00m 08s

Summary and Conclusion

Evaluation


```

In [12]: # Get the best model
best_model = tuner.get_best_models()[0]

# Training the best model
history = best_model.fit(X_train, y_train, epochs=50, validation_data=(X_

# Evaluation
loss, accuracy = best_model.evaluate(X_test, y_test)
print(f"Accuracy on test set: {accuracy*100:.2f}%")

Epoch 45/50
15/15 [=====] - 0s 2ms/step - loss: 0.0331 - a
ccuracy: 0.9890 - val_loss: 0.1313 - val_accuracy: 0.9474
Epoch 46/50
15/15 [=====] - 0s 2ms/step - loss: 0.0568 - a
ccuracy: 0.9758 - val_loss: 0.1159 - val_accuracy: 0.9561
Epoch 47/50
15/15 [=====] - 0s 2ms/step - loss: 0.0294 - a
ccuracy: 0.9956 - val_loss: 0.1317 - val_accuracy: 0.9649
Epoch 48/50
15/15 [=====] - 0s 2ms/step - loss: 0.0269 - a
ccuracy: 0.9890 - val_loss: 0.1156 - val_accuracy: 0.9649
Epoch 49/50
15/15 [=====] - 0s 2ms/step - loss: 0.0309 - a
ccuracy: 0.9912 - val_loss: 0.1144 - val_accuracy: 0.9649
Epoch 50/50
15/15 [=====] - 0s 2ms/step - loss: 0.0263 - a
ccuracy: 0.9912 - val_loss: 0.1083 - val_accuracy: 0.9649
4/4 [=====] - 0s 922us/step - loss: 0.1083 - a
ccuracy: 0.9649
Accuracy on test set: 96.49%

```

```
In [13]: from collections import Counter
from sklearn.metrics import confusion_matrix
import itertools
# From the SKLEARN website, to view a confusion matrix
def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

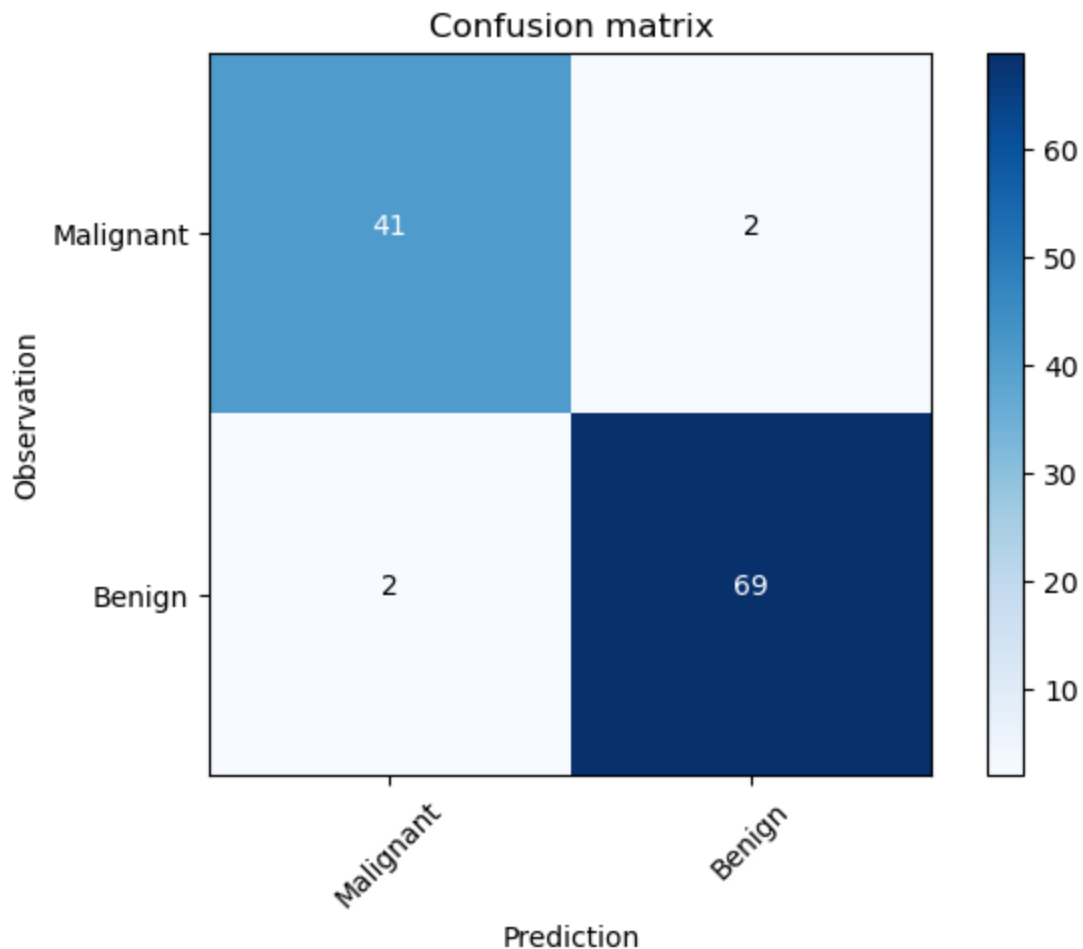
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, cm[i, j],
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('Observation')
    plt.xlabel('Prediction')
```

```
In [14]: # Confusion Matrix
from sklearn.metrics import confusion_matrix, classification_report
y_pred = (best_model.predict(X_test) > 0.5).astype("int32")
cm = confusion_matrix(y_test, y_pred)
class_names = ['Malignant', 'Benign']
plot_confusion_matrix(cm, class_names)
plt.show()
```

4/4 [=====] - 0s 738us/step



```
In [15]: print(classification_report(y_test, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.95 | 0.95 | 0.95 | 43 |
| 1 | 0.97 | 0.97 | 0.97 | 71 |
| accuracy | | | 0.96 | 114 |
| macro avg | 0.96 | 0.96 | 0.96 | 114 |
| weighted avg | 0.96 | 0.96 | 0.96 | 114 |

Explain the confusion matrix with your own words


1. The model predicted 41 cases as "Malignant" and they were actually "Malignant".

2. The model predicted 2 cases as "Benign" when they were actually "Malignant".
3. The model predicted 2 cases as "Malignant" when they were actually "Benign".
4. The model predicted 69 cases as "Benign" and they were actually "Benign".

Some observations:

- The model seems to be performing quite well as the majority of predictions fall on the diagonal, which represents correct predictions.
- The errors are balanced, with the model misclassifying 2 cases for both false positives and false negatives.

```
In [16]: # Check for overfitting in the history object
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')
plt.show()
```



Overfitting Observations

- **High Training Accuracy:** The blue line, which represents the training accuracy, is consistently high, nearly reaching 1.00 (or 100%) for most epochs. This suggests that the model has learned the training data very well.
- **Validation Accuracy Fluctuations:** The orange line, representing validation accuracy, shows more fluctuation compared to the training accuracy. There's a noticeable dip in the middle epochs and then it rises again towards the later epochs. This fluctuation suggests

that the model might be experiencing some variability in how well it generalizes to unseen data.

- **Divergence between Training and Validation:** Around the middle epochs (approximately epochs 20-35), there's a clear gap between training and validation accuracy. This gap suggests that the model might be overfitting the training data during these epochs, as it performs exceptionally well on the training data but not as well on the validation data.
- **Convergence in Later Epochs:** Towards the later epochs (after 40), the validation accuracy seems to improve and get closer to the training accuracy. This convergence indicates that the model's generalization to unseen data has improved in these epochs.

In []: